



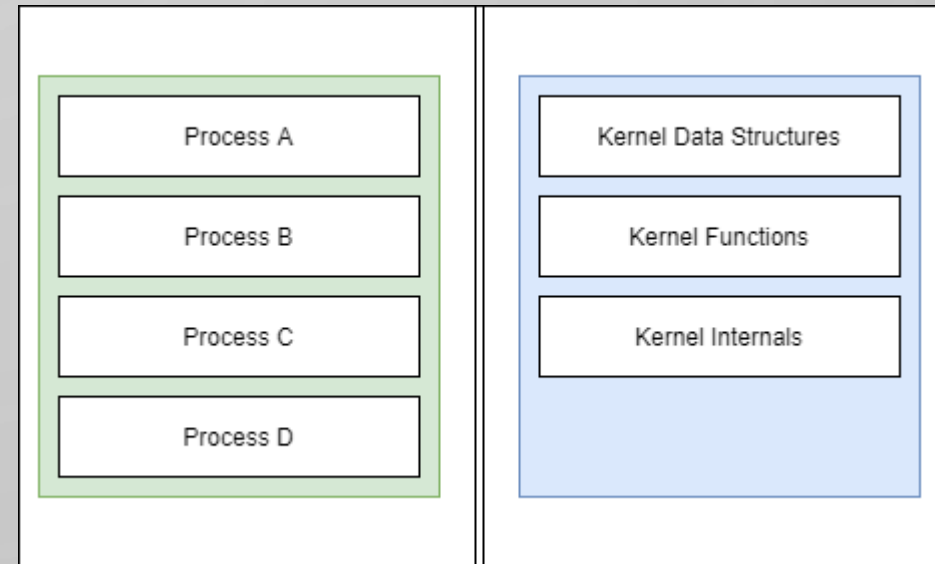
SYSTEM CALLS AND KERNEL SERVICES

CPE / CSC 159: OPERATING SYSTEM PRAGMATICS

GREG CRIST (CRIST@CSUS.EDU)

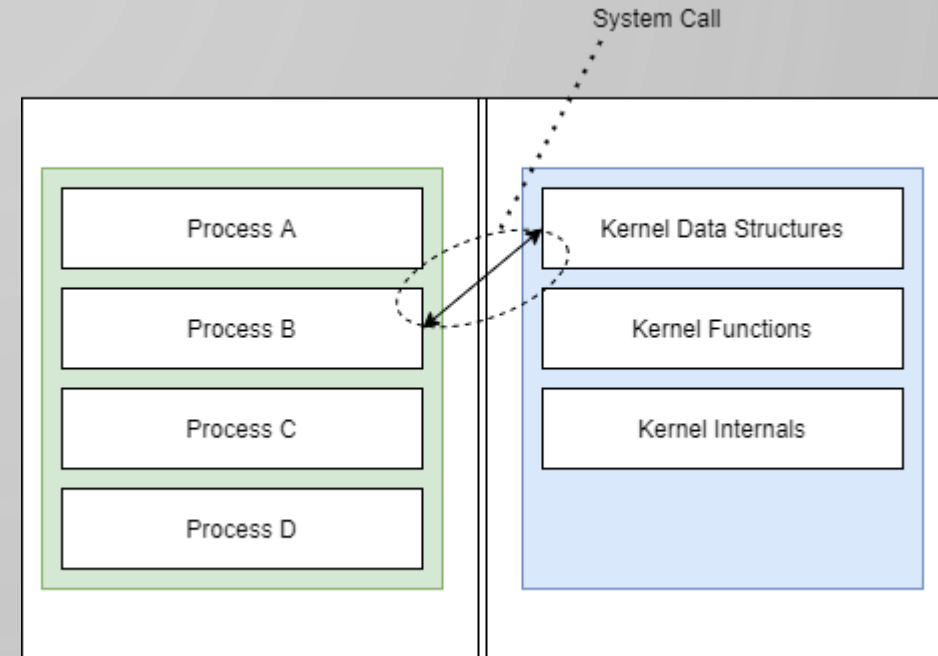
THE KERNEL/USER-SPACE BARRIER

- Maintains a separation between kernel-space and user-space
- The kernel has visibility into the operation of the processes that run, but processes shouldn't have visibility into the kernel operation



THE KERNEL/USER-SPACE BARRIER

- All processes interact with the kernel services through the use of **system calls**.



WHAT ARE KERNEL SERVICES?

- The Kernel maintains a set of services that it provides, such as:
 - Process management
 - Memory management
 - System information
 - Communication
 - Hardware interfacing / drivers

EXAMPLES OF SYSTEM CALLS

- Process Management
 - Creating a new process
 - Exiting a process
 - Putting a process to “sleep”
- Memory Management
 - Memory allocation (ex: malloc, free)
- Communication
 - Interprocess communication
 - Networking
- File management
 - Open / close a file
 - Read from / write to a file
- Retrieving kernel/system information
 - Obtaining the system time
 - Obtaining the running process id
- Hardware Interfacing
 - Input/Output operations

HIGH-LEVEL OVERVIEW OF SYSTEM CALLS

- System calls provide mechanisms to:
 - Exchange data between kernel/runtime processes
 - Enter the kernel context to access the kernel service
 - Return to the user process
- How might this be performed?
 - How to enter the kernel context?
 - When in the kernel context, how to exchange data?
 - How to return to the user process?

SYSTEM CALLS: ENTERING KERNEL CONTEXT

- Kernel context can be triggered via interrupts
 - Hardware driven interrupts (such as the timer interrupt) enter the kernel for processing
 - Software driven interrupts can similarly be used to enter the kernel context
- Software instructions may be executed to enter kernel context as well

SYSTEM CALLS: DATA EXCHANGE

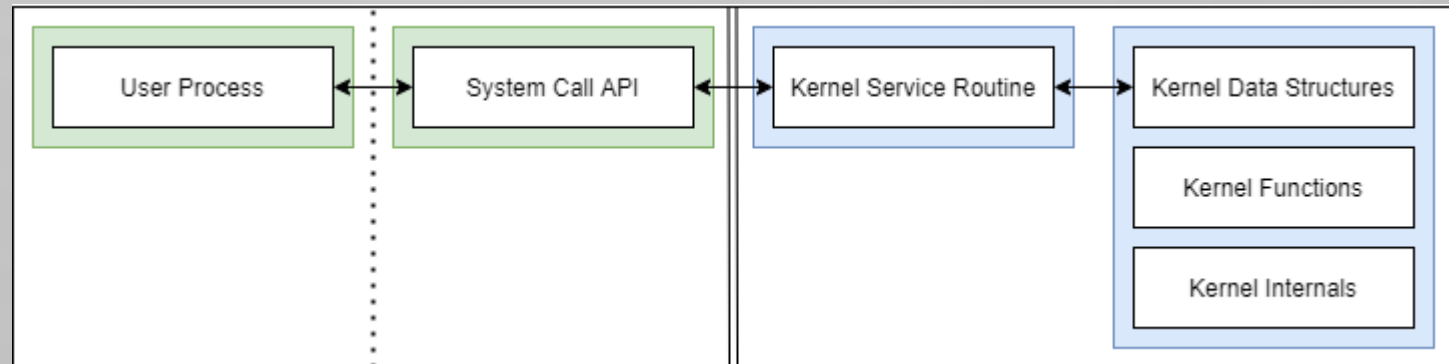
- When an interrupt is triggered, the CPU state is saved for the interrupt service routine to run
- This is leveraged to exchange data using CPU registers
 - Data can be sent from user space to kernel space
 - Data can be received from kernel space to user space

SYSTEM CALLS: RETURNING TO USER SPACE

- Just as a hardware interrupt is triggered and processed, a software interrupt is triggered and processed
- When the interrupt servicing is complete, the user process may be loaded and continue executing, completing the system call

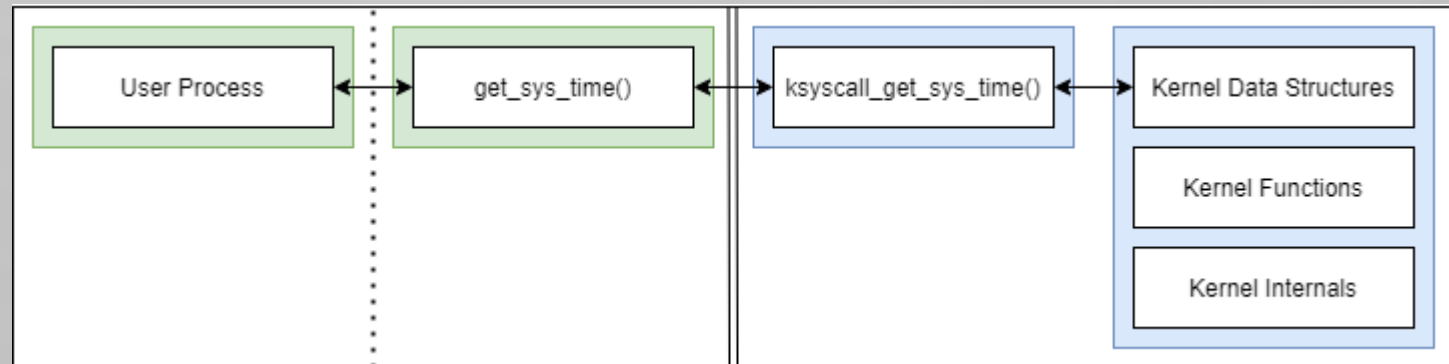
WHAT MAKES UP A SYSTEM CALL?

- System calls are made up of various components:
 - The system call API (function called from user space)
 - The interrupt to trigger the system call
 - Interrupt service routine to enter kernel context
 - System call handler



EXAMPLE SYSTEM CALL: GET_SYS_TIME()

- User process calls an API (function) to retrieve the current system time
 - `get_sys_time()` function triggers the system call interrupt
 - ISR results in kernel context being entered (`kernel_run`) and is handled
 - System call is dispatching to the specific system call handler (`ksyscall_get_sys_time()`)
 - Process scheduler loads the process and resumes to user context
 - System call API completes and returns to the calling process



GET_SYS_TIME() – USER CONTEXT

- User process calls the system call API that:
 - Sets registers to indicate to the kernel the system call to perform
 - Triggers an interrupt to enter kernel context
 - When the interrupt returns, sets data and then finally returns to the caller

```
void user_proc(void) {  
    int current_time;  
    current_time = get_sys_time();  
    printf("System time is: %d\n", current_time);  
}
```

```
int get_sys_time() {  
    int time;  
    asm("movl %1, %%eax;"  
        "int $0x80;"  
        "movl %%ebx, %0;"  
        : "=g" (time)  
        : "g" (SYSCALL_GET_SYS_TIME)  
        : "eax", "ebx");  
    return time;  
}
```

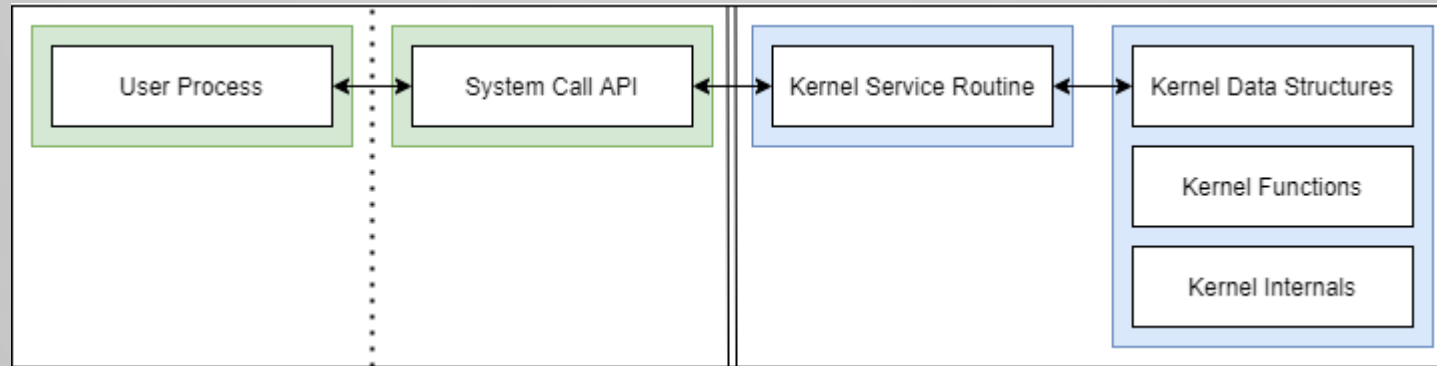
GET_SYS_TIME(): KERNEL CONTEXT

- In the kernel:
 - Syscall Interrupt Service Routine decodes the system call
 - Dispatches to the kernel service routine to process
 - Continues kernel run loop
 - Process is loaded / restored -> CPU registers are set

```
void ksr_syscall(void) {  
    int syscall = pcb[active_pid].trapframe_p->eax;  
    if (syscall == SYSCALL_GET_SYS_TIME) {  
        ksyscall_get_sys_time();  
    }  
}
```

```
void ksyscall_get_sys_time() {  
    pcb[active_pid].trapframe_p->ebx = system_time;  
}
```

REVISITING THE SYSTEM CALL CONTEXT SWITCH



- All system calls trigger the context switch from user context to kernel context and then back to user context
- Only mechanism for data exchange is via CPU registers
 - Consider: what can you pass through registers?

SYSTEM CALLS: BEYOND 80386

- Different CPU architectures provide different mechanisms; we are looking at 80386 (subset of larger x86 instruction set)
 - x86: use a software defined interrupt to enter kernel context
 - IA-32 (Intel Pentium) introduced the SYSENTER instruction to trigger a system call
 - SYSEXIT instruction to return to the caller
 - x64 introduced the SYSCALL instruction to trigger a system call
 - SYSRET instruction to return to the caller
 - ARM64: uses the SVC instruction (supervisor call) to trigger a system call
- Dedicated instructions are used to for speed/simplicity
 - Interrupts are generally considered expensive
 - Dedicated instructions can perform additional work versus needing to write additional code